

TAI NGUYEN PHU

+84 945 409 269 tainguyenphu2502@gmail.com linkedin.com/in/yuitec/ github.com/YuITC/

Aspiring AI Engineer specializing in LLM-powered systems, with a focus on Retrieval-Augmented Generation, agentic workflows, and model fine-tuning. Seeking to contribute to AI-driven product teams.

EDUCATION

Master of Computer Science	University of Information Technology - VNUHCM Dec 2025 - Present
Bachelor of Computer Science GPA: 3.5/4.0	University of Information Technology - VNUHCM Sep 2022 - Sep 2025

EXPERIENCE

AI Engineer Intern — CoverGo Insurtech	Oct 2025 - Jan 2026
- Researched DSPy-based LLM optimization techniques to improve consistency in the product pipeline - Conducted prompt engineering and Gemini integration experiments to enhance product workflows - Built a RAG-based Q&A web application leveraging knowledge acquired during the internship	

PROJECTS

Metacognitive RAG System with Adaptive Strategy Selection	Jan 2026 - Mar 2026
- Built a multi-step AI research agent using LangGraph, enabling accurate multi-source question answering, achieving 59% exact match on benchmark datasets, significantly outperforming baseline methods - Designed a self-correction system that classifies answer failures and auto-applies targeted fixes, reducing incorrect responses by 46–66% prior to final output - Integrated hybrid search using Qdrant and BM25, improving document retrieval relevance - Developed an adaptive strategy selector using Thompson Sampling (auto-learning optimization), enabling the system to learn the best retrieval settings per query type over time	
Novel-Verse AI	Nov 2025 - Jan 2026
- Built a full-stack Next.js AI-first novel reading web application with DeepSeek and OpenAI embeddings - Integrated a RAG system using OpenAI embeddings and PostgreSQL vector search, enabling accurate AI-powered Q&A grounded in story content - Built a multi-layer crawler with automatic fallback mechanisms, bypassing anti-bot protections for reliable data - Developed real-time SSE streaming for crawl and translation progress, providing live status updates for long-running AI tasks	
Vietnamese Legal Document Retrieval	Mar 2025 - May 2025
- Fine-tuned Sentence-BERT on over 100K Vietnamese legal documents - Evaluated retrieval model via MTEB benchmark, achieving 79% recall, greatly outperforming baseline model - Built FAISS vector index for fast legal document search system, enabling sub-second legal search - Deployed end-to-end legal document search app with Docker and Gradio on Hugging Face Spaces	

SKILLS

Programming	Python, C++, TypeScript, SQL, REST API, FastAPI, OOP, Git
AI/LLM	RAG, LangChain, LangGraph, LoRA, QLoRA, DSPy, Prompt Engineering, AI Agents
ML & Data	Qdrant, FAISS, Transformers, Pandas, NumPy, PyTorch, Scikit-learn, OpenCV
MLOps	Docker, AWS, MLflow, Weights & Biases, GitHub Actions
Soft Skills	English (IELTS 7.0), Research, Presentation, Teamwork, Project Management

CERTIFICATIONS

NVIDIA — Applications of AI for Anomaly Detection	Mar 2026
IELTS Academic (7.0 Overall)	Oct 2025
Deep Learning Specialization	Aug 2025
Machine Learning Specialization	Jul 2025